



## **Zeyobron Bigdata Analytics**

### **Project phase-1**

**Name : Navaneetha Krishnan**

**Mobile no : 9600730899**

**Project Reporting : Mr - Sai Adithya**

### **Contents**

- Verify you have customer\_total table
- Create customer\_src table like customer\_total
- Insert into customer\_src from customer\_total where id>300 and id<350
- Create a password file your edge Node == (echo -n cloudera>passfile)
- Create a directory in Edge Node
- Go inside that directory
- Create sqoop job for customer\_src table with id last value 0 along with psw\_file
- Execute the Job
- Verify whether data got imported
- AVSC file should have generated locally-- Copy that to HDFS
- Go inside Hive shell and create database prodb
- Now create Hive table on top of raw data with AVSC FILE as reference



## Project phase-1

### 1) Verify you have customer\_total table

Login to MySQL [ [mysql -uroot -pcloudera](#) ]

View all databases [ [show databases;](#) ]

```
mysql> show databases;
+-----+
| Database |
+-----+
| information_schema |
| addl      |
| cm        |
| dataaa    |
| datap     |
| datap1    |
| db1       |
| firehose  |
| hue       |
| inc       |
| incjob    |
| metastore |
| mysql     |
| nav       |
| navms     |
| oozie     |
| prac1     |
| prac2     |
| prodb     |
| retail_db |
+-----+
```

Select prodb [ [use prodb;](#) ]

View all tables [ [show tables;](#) ]

```
mysql> use prodb;
Reading table information for completion of table and column names
You can turn off this feature to get a quicker startup with -A

Database changed
mysql> show tables;
+-----+
| Tables_in_prodb |
+-----+
| customer_total  |
+-----+
1 row in set (0.00 sec)
```

## 2) Create customer\_src table like customer\_total

Command to create customer\_src table

```
[ create table customer_src(id int(10),username varchar(100),sub_port varchar(100),host
varchar(100),date_time varchar(100),hit_count_var_1 varchar(100),hit_count_var_2
varchar(100),hit_count_var_3 varchar(100),timezone varchar(100),method
varchar(100),procedure varchar(100),value varchar(100),sub_product
varchar(100),web_info varchar(100),status_code varchar(100)); ]
```

```
mysql> desc customer_total;
```

Field	Type	Null	Key	Default	Extra
id	int(10)	YES		NULL	
username	varchar(100)	YES		NULL	
sub_port	varchar(100)	YES		NULL	
host	varchar(100)	YES		NULL	
date_time	varchar(100)	YES		NULL	
hit_count_val_1	varchar(100)	YES		NULL	
hit_count_val_2	varchar(100)	YES		NULL	
hit_count_val_3	varchar(100)	YES		NULL	
timezone	varchar(100)	YES		NULL	
method	varchar(100)	YES		NULL	
procedure	varchar(100)	YES		NULL	
value	varchar(100)	YES		NULL	
sub_product	varchar(100)	YES		NULL	
web_info	varchar(100)	YES		NULL	
status_code	varchar(100)	YES		NULL	

View all tables [ [show tables;](#) ]

```
mysql> show tables;
+-----+
| Tables_in_prodb |
+-----+
| customer_src     |
| customer_total   |
+-----+
2 rows in set (0.00 sec)
```

## 3) Insert into customer\_src from customer\_total where id>300 and id<350

View all tables [ [insert into customer\\_src select \\* from customer\\_total where id>300 and id<350;](#) ]



#### 4) Create a password file your edge Node == (echo -n cloudera>passfile)

Create edge node pass file [ `echo -n Adityago8>/home/itv004084/apassfile` ]

#### 5) Create a directory in Edge Node

Directory in edge node [ `mkdir pharsedir` ]

#### 6) Go inside that directory

Directory inside edge node [ `cd pharsedir` ]

#### 7) Create sqoop job for customer\_src table with id last value 0 along with password file

sqoop job delete command [ `sqoop job --delete pjob` ]

sqoop job create [ `sqoop job --create pjob -- import --connect jdbc:mysql://localhost/datap --username root --password-file file:///home/cloudera/pfile --m 1 --table intab --target-dir /user/cloudera/datap --incremental append --check-column id --last-value 0 --as-avrodatafile` ]

```
[itv004084@g02 ~]$ sqoop job --create mobilejob -- import --connect jdbc:mysql://zeyod2.czgbk54pezaa.ap-south-1.rds.amazonaws.com/itv004084 --username root --password-file file:///home/itv004084/pfile --m 1 --table itv004084 --target-dir /user/itv004084/datap --incremental append --check-column id --last-value 0
Warning: /opt/sqoop/../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /opt/sqoop/../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/opt/hadoop-3.3.0/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/opt/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
2022-11-03 02:53:34,215 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
[itv004084@g02 ~]$
```

#### 8) Execute the Job

sqoop job execute command [ `sqoop job --exec pjob` ]

```
[itv004084@g02 ~]$ sqoop job --list
Warning: /opt/sqoop/../hcatalog does not exist! HCatalog jobs will fail.
Please set $HCAT_HOME to the root of your HCatalog installation.
Warning: /opt/sqoop/../accumulo does not exist! Accumulo imports will fail.
Please set $ACCUMULO_HOME to the root of your Accumulo installation.
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/opt/hadoop-3.3.0/share/hadoop/common/lib/slf4j-log4j12-1.7.25.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/opt/apache-hive-3.1.2-bin/lib/log4j-slf4j-impl-2.10.0.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
2022-11-03 02:56:12,229 INFO sqoop.Sqoop: Running Sqoop version: 1.4.7
Available jobs:
  incjob
  incjob
  incjob
  incjobs
  itv_job
  itvjjob
  mobilejob
  naveenjob
  ncjob
  newj
  pjob
  pjob
  pjob
[itv004084@g02 ~]$
```



## 9) Verify whether data got imported

Data importing command [ `hadoop fs -ls /user/itv004084/customer_data` ]

```
drwx----- - itv004084 supergroup          0 2022-10-16 02:17 /user/itv004084/.Trash
drwxr-xr-x - itv004084 supergroup          0 2022-11-03 02:59 /user/itv004084/_sqoop
drwxr-xr-x - itv004084 supergroup          0 2022-10-17 13:41 /user/itv004084/aavdirr
drwxr-xr-x - itv004084 supergroup          0 2022-10-03 02:09 /user/itv004084/adata
drwxr-xr-x - itv004084 supergroup          0 2022-10-30 03:38 /user/itv004084/adir
-rw-r--r--  3 itv004084 supergroup        508 2022-10-30 03:28 /user/itv004084/atab.avsc
drwxr-xr-x - itv004084 supergroup          0 2022-10-17 11:39 /user/itv004084/avdir
drwxr-xr-x - itv004084 supergroup          0 2022-10-17 13:37 /user/itv004084/avdirr
drwxr-xr-x - itv004084 supergroup          0 2022-09-24 20:39 /user/itv004084/avrodata
drwxr-xr-x - itv004084 supergroup          0 2022-10-30 03:49 /user/itv004084/avscdirpro
drwxr-xr-x - itv004084 supergroup          0 2022-09-22 00:58 /user/itv004084/chdata
drwxr-xr-x - itv004084 supergroup          0 2022-09-25 11:46 /user/itv004084/chennaidir
drwxr-xr-x - itv004084 supergroup          0 2022-10-16 02:43 /user/itv004084/confirm
drwxr-xr-x - itv004084 supergroup          0 2022-10-30 03:47 /user/itv004084/customer_stage_loc
-rw-r--r--  3 itv004084 supergroup        13 2022-10-16 02:20 /user/itv004084/data.csv
drwxr-xr-x - itv004084 supergroup          0 2022-09-22 00:51 /user/itv004084/datadir
drwxr-xr-x - itv004084 supergroup          0 2022-11-03 02:59 /user/itv004084/datap
drwxr-xr-x - itv004084 supergroup          0 2022-10-16 00:54 /user/itv004084/ddir
drwxr-xr-x - itv004084 supergroup          0 2022-10-06 12:51 /user/itv004084/diinimport
drwxr-xr-x - itv004084 supergroup          0 2022-10-04 01:38 /user/itv004084/dinimport
drwxr-xr-x - itv004084 supergroup          0 2022-10-08 03:11 /user/itv004084/djinr
drwxr-xr-x - itv004084 supergroup          0 2022-10-31 10:51 /user/itv004084/dyndir
drwxr-xr-x - itv004084 supergroup          0 2022-10-28 02:53 /user/itv004084/edir
drwxr-xr-x - itv004084 supergroup          0 2022-10-16 02:38 /user/itv004084/exdir
drwxr-xr-x - itv004084 supergroup          0 2022-09-19 00:15 /user/itv004084/h_zzdir
drwxr-xr-x - itv004084 supergroup          0 2022-09-26 12:59 /user/itv004084/indir
drwxr-xr-x - itv004084 supergroup          0 2022-10-06 14:03 /user/itv004084/jobdir
drwxr-xr-x - itv004084 supergroup          0 2022-09-26 12:18 /user/itv004084/joindir
drwxr-xr-x - itv004084 supergroup          0 2022-09-26 16:09 /user/itv004084/naveen
drwxr-xr-x - itv004084 supergroup          0 2022-09-22 00:26 /user/itv004084/newimp
drwxr-xr-x - itv004084 supergroup          0 2022-09-26 15:38 /user/itv004084/pardir
drwxr-xr-x - itv004084 supergroup          0 2022-10-02 00:52 /user/itv004084/pdata
drwxr-xr-x - itv004084 supergroup          0 2022-10-22 01:03 /user/itv004084/pdir
drwxr-xr-x - itv004084 supergroup          0 2022-10-16 01:05 /user/itv004084/rdir
drwxr-xr-x - itv004084 supergroup          0 2022-10-22 23:41 /user/itv004084/sdir
drwxr-xr-x - itv004084 supergroup          0 2022-10-22 23:41 /user/itv004084/sidir
drwxr-xr-x - itv004084 supergroup          0 2022-10-27 12:38 /user/itv004084/srcd
drwxr-xr-x - itv004084 supergroup          0 2022-10-30 03:41 /user/itv004084/stdir
drwxr-xr-x - itv004084 supergroup          0 2022-10-31 10:59 /user/itv004084/subdyndir
drwxr-xr-x - itv004084 supergroup          0 2022-10-31 10:54 /user/itv004084/subsrcdir
drwxr-xr-x - itv004084 supergroup          0 2022-10-27 12:38 /user/itv004084/tard
drwxr-xr-x - itv004084 supergroup          0 2022-10-31 10:31 /user/itv004084/warehouse
-rw-r--r--  3 itv004084 supergroup          9 2022-09-19 01:29 /user/itv004084/zfile
[itv004084@g02 ~]$
```

## 10) AVSC file should have generated locally-- Copy that to HDFS

Create one directory command [ `hadoop fs -mkdir /user/cloudera/projectavsc` ]

Move file edge node command [ `hadoop fs -put /home/cloudera/phase1/customer_src.avsc /user/cloudera/projectavsc` ]

```
[cloudera@quickstart phase1]$ ls -l
total 44
-rw-rw-r-- 1 cloudera cloudera 2280 Oct 29 13:24 customer_src.avsc
-rw-rw-r-- 1 cloudera cloudera 38088 Oct 29 13:24 customer_src.java
[cloudera@quickstart phase1]$
```





## 11) Go inside Hive shell and create database prodb;

Hive shell command [ `hive` ], [ `create database prodb;` ], [ `use prodb;` ]

```
[cloudera@quickstart ~]$ hadoop fs -ls
Found 25 items
drwxr-xr-x - cloudera cloudera 0 2022-10-29 13:26 _sqoop
drwxr-xr-x - cloudera cloudera 0 2022-10-06 12:14 adata
drwxr-xr-x - cloudera cloudera 0 2022-10-29 10:39 adir
-rw-r--r-- 1 cloudera cloudera 508 2022-10-29 10:19 atab.avsc
drwxr-xr-x - cloudera cloudera 0 2022-09-17 23:09 chdata
drwxr-xr-x - cloudera cloudera 0 2022-09-23 21:20 chennaidir
drwxr-xr-x - cloudera cloudera 0 2022-10-29 13:26 customer_data
-rw-r--r-- 1 cloudera cloudera 29156 2022-09-16 22:59 data.avro
drwxr-xr-x - cloudera cloudera 0 2022-09-17 21:59 datadir
drwxr-xr-x - cloudera cloudera 0 2022-09-10 22:10 datar
```

## 12) Now create Hive table on top of raw data with AVSC FILE as reference

Hive raw data command [ `create table customer_src row format serde 'org.apache.hadoop.hive.serde2.avro.AvroSerDe' STORED AS AVRO LOCATION '/user/cloudera/customer_data' TBLPROPERTIES ('avro.schema.url'='/user/cloudera/projectavro/customer_src.avsc');` ]

307	beautifulbear	06865	664.687.377.602	12/Sep/2012:18:09:29	-0500	18	09	29	-0500	GET	/download/download3.zip	0	/download	M
308	beautifulbear	01759	668.686.42.61	28/Dec/2011:00:32:33	-0500	00	32	33	-0500	GET	/partners	0	/partners	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT 5.1; .NET CLR 2.0.50727; .NET CLR 3.0.04506.30; InfoPa
309	beautifulbear	01701	14.323.74.653	13/Sep/2012:02:29:30	-0500	02	29	30	-0500	GET	/demo	0	-	Jakarta Commons-HttpClie
310	beautifulbear	43137	52.43.672.362	04/Jan/2012:12:12:06	-0500	12	12	06	-0500	GET	/docs/doc2.pdf	0	/docs/doc2.pdf	Mozilla/5.0 (Windows; U; Windows NT 5.1; de; rv:1.9.2.3) Gecko/20100401 Firefox/3.6.3 (.NET CLR 3.5
311	beautifulbear	04126	53.378.683.602	14/Sep/2012:13:11:32	-0500	13	11	32	-0500	GET	/feeds/press	0	/feeds/press	Mozilla/5.0 (Windows; U; Windows NT 6.0; fr; rv:1.9.2.3) Gecko/20100401 Firefox/3.6.3 (.NET CLR 3.5
312	beautifulbear	32800	75.622.36.380	05/Jan/2012:09:12:28	-0500	09	12	28	-0500	GET	/products	0	/products	Mozilla/5.0 (Windows; U; Windows NT 5.1; en-US; rv:1.9.1.5) Gecko/20091102 Firefox/3.5.5 GTB5
313	beautifulbear	02621	14.610.30.67	18/Sep/2012:08:54:30	-0500	08	54	30	-0500	GET	/product/product3	0	-	Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 5.2; Trident/4.0; .NET CLR 1.1.4322; .NET CLR 2.0.5072
314	beautifulbear	80516	680.330.387.7	05/Jan/2012:12:48:38	-0500	12	48	38	-0500	GET	/product/product3	0	/product/product	Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 5.1; Trident/4.0; .NET CLR 1.1.4322)
315	beautifulbear	94600	361.631.17.30	20/Sep/2012:23:24:59	-0500	23	24	59	-0500	GET	/demo	0	/demo	Jakarta Commons-HttpClie
316	beautifulbear	08102	322.6.648.325	05/Jan/2012:14:05:22	-0500	14	05	22	-0500	GET	/download/download3.zip	0	/download/downlo	Mozilla/4.0 (compatible; MSIE 5.00; Windows 98) 100
317	beautifulbear	04731	87.86.81.40	21/Sep/2012:05:54:30	-0500	05	54	30	-0500	GET	/feeds/press	0	-	Mozilla/5.0 (Win
318	beautifulbear	06435	30.672.8.336	06/Jan/2012:20:01:36	-0500	20	01	36	-0500	GET	/home	0	/home	Jakarta Commons-HttpClie
319	beautifulbear	06597	367.34.686.631	22/Sep/2012:23:09:30	-0500	23	09	30	-0500	GET	/demo	0	-	Jakarta Commons-HttpClie
320	beautifulbear	90110	322.6.648.325	06/Jan/2012:20:30:54	-0500	20	30	54	-0500	GET	/download/download3.zip	0	/download/downlo	Mozilla/4.0 (compatible; MSIE 5.00; Windows 98) 100
321	beautifulbear	61561	13.626.82.338	24/Sep/2012:19:11:38	-0500	19	11	38	-0500	GET	/product/product3	0	/product/product	Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 5.1; Trident/4.0; Mozilla/4.0 (compatible; MSIE 6.0; W
322	beautifulbear	01794	650.47.72.628	06/Jan/2012:21:22:30	-0500	21	22	30	-0500	GET	/misc	0	-	Python-urllib/2.5
323	beautifulbear	60546	41.631.382.601	25/Sep/2012:22:40:00	-0500	22	40	00	-0500	GET	/feeds/press	0	-	Apple-PubSub/65.
324	beautifulbear	04825	322.6.648.325	07/Jan/2012:06:32:49	-0500	06	32	49	-0500	GET	/download/download3.zip	0	/download/downlo	Mozilla/4.0 (compatible; MSIE 5.00; Windows 98) 100
325	beautifulbear	61449	361.54.46.628	27/Sep/2012:12:07:24	-0500	12	07	24	-0500	GET	/partners	0	/partners	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT 6.0; Mozilla/4.0 (compatible; MSIE 6.0; Windows NT 5.1
326	beautifulbear	40 53	322.6.648.325	07/Jan/2012:09:04:53	-0500	09	04	53	-0500	GET	/download/download3.zip	0	/download/downlo	Mozilla/4.0 (compatible; MSIE 5.00; Windows 98) 100
327	beautifulbear	09264	628.44.58.626	30/Oct/2011:07:26:30	-0500	07	26	30	-0500	GET	/feeds/press	0	-	Apple-PubSub/65.
328	beautifulbear	21605	322.6.648.325	08/Jan/2012:05:18:12	-0500	05	18	12	-0500	GET	/download/download3.zip	0	/download/downlo	